

class07

Sharlene Yang


Clustering

First let's make up some data to cluster so we can get a feel for these methods and how to work with them.

We can use the `rnorm()` function to get random numbers from a normal distribution around a given `mean`.

```
hist(rnorm(5000, mean = 3))
```

Histogram of `rnorm(5000, mean = 3)`



Let's get 30 points with a mean of 3.

```
tmp <- c(rnorm(30, mean= 3), rnorm(30, mean = -3))
tmp
```

```
[1] 1.854139 2.662916 3.763941 3.276596 1.383767 3.378643 4.011393
[8] 2.270529 3.187053 3.166477 3.005245 2.688300 4.234360 4.141733
[15] 2.073393 3.349752 5.098710 2.722750 3.012937 3.495058 4.243105
[22] 4.414468 3.512653 4.469846 2.233505 1.943427 3.134889 3.198945
[29] 3.738780 2.869347 -2.456904 -4.087236 -4.520312 -3.938465 -2.806694
[36] -1.800357 -3.836458 -1.917171 -3.656225 -4.677399 -4.149475 -4.196859
[43] -1.691379 -1.851419 -2.628294 -2.269254 -4.471422 -3.615363 -2.097901
[50] -5.610092 -2.528510 -3.066923 -2.953418 -4.621749 -2.890803 -1.546576
[57] -3.735499 -2.620206 -2.287271 -4.369600
```


Put two of these together:

```
x <- cbind(x=tmp, y= rev(tmp))
x
```

```
      x      y
[1,] 1.854139 -4.369600
[2,] 2.662916 -2.287271
[3,] 3.763941 -2.620206
[4,] 3.276596 -3.735499
[5,] 1.383767 -1.546576
[6,] 3.378643 -2.890803
[7,] 4.011393 -4.621749
[8,] 2.270529 -2.953418
[9,] 3.187053 -3.066923
[10,] 3.166477 -2.528510
[11,] 3.005245 -5.610092
[12,] 2.688300 -2.097901
[13,] 4.234360 -3.615363
[14,] 4.141733 -4.471422
[15,] 2.073393 -2.269254
[16,] 3.349752 -2.628294
[17,] 5.098710 -1.851419
[18,] 2.722750 -1.691379
[19,] 3.012937 -4.196859
[20,] 3.495058 -4.149475
[21,] 4.243105 -4.677399
[22,] 4.414468 -3.656225
```

```
[23,] 3.512653 -1.917171
[24,] 4.469846 -3.836458
[25,] 2.233505 -1.800357
[26,] 1.943427 -2.806694
[27,] 3.134889 -3.938465
[28,] 3.198945 -4.520312
[29,] 3.738780 -4.087236
[30,] 2.869347 -2.456904
[31,] -2.456904 2.869347
[32,] -4.087236 3.738780
[33,] -4.520312 3.198945
[34,] -3.938465 3.134889
[35,] -2.806694 1.943427
[36,] -1.800357 2.233505
[37,] -3.836458 4.469846
[38,] -1.917171 3.512653
[39,] -3.656225 4.414468
[40,] -4.677399 4.243105
[41,] -4.149475 3.495058
[42,] -4.196859 3.012937
[43,] -1.691379 2.722750
[44,] -1.851419 5.098710
[45,] -2.628294 3.349752
[46,] -2.269254 2.073393
[47,] -4.471422 4.141733
[48,] -3.615363 4.234360
[49,] -2.097901 2.688300
[50,] -5.610092 3.005245
[51,] -2.528510 3.166477
[52,] -3.066923 3.187053
[53,] -2.953418 2.270529
[54,] -4.621749 4.011393
[55,] -2.890803 3.378643
[56,] -1.546576 1.383767
[57,] -3.735499 3.276596
[58,] -2.620206 3.763941
[59,] -2.287271 2.662916
[60,] -4.369600 1.854139
```

```
plot(x)
```



K-means clustering.

Very popular clustering method that we can use with the `kmeans()` function in base R.

```
km <- kmeans(x, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.229975	3.217889
2	3.217889	-3.229975

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 56.31718 56.31718
(between_SS / total_SS = 91.7 %)
```

Available components:

[1]	"cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6]	"betweenss"	"size"	"iter"	"ifault"	


```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
km$center
```


	x	y
1	-3.229975	3.217889
2	3.217889	-3.229975

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch = 15)
```



Q. Let's cluster into 3 groups or some x data and make a plot

```
km <- kmeans(x, centers = 4)
plot(x, col=km$cluster)
```



Hierarchical Clustering

We can use the `hclust()` for Hierarchical Clustering Unlike `kmeans()`, where we could just pass in our data as input, we need to give `hclust()` a “distance matrix”.

We will use the `dist()` function to start with.

```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```




I

[

Y

I



Principal Component Analysis (PCA)

Class 7 lab

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

There are 17 rows and 5 columns.


```
View(x)
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
tail(x)
```

	X	England	Wales	Scotland	N.Ireland
12	Fresh_fruit	1102	1137	957	674
13	Cereals	1472	1582	1462	1494
14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

To fix the name of the row, add `row.name = 1` to the `read.csv`

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names= 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334

Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Checking the dimensions again


```
dim(x)
```

```
[1] 17  4
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer the adding `row.names` because if you use the `-1` method and continuously run it, it will remove part of the table, which isn’t what we want.


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


Changing the “besides” argument will change it to stacked bars.

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



For the code, it graphs the countries on the x-axis or the y-axis. It compares the amount of products consumed/ used in each country. For example, row 1 column 2, it shows England on the y-axis and Wales on the x-axis. If a given point lies on the diagonal for a given plot, it means that both countries have used about the same amount of the product.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Main difference is that it has the distinct blue dot near the 1000 mark.

The main PCA function in base R is called `prcomp()` it expects the transpose of our data.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names  
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class  
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
plot(pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"), pch= 16)
```

